



PRISM NODE

The AI Voice Buyer's Guide

What to look for, what to refuse to pay for, and the seven questions every vendor dodges.

By the team at TrainYourAgent · Tampa Bay, FL · 2026 edition

cal.com/trainyouragent/30min

Table of contents

- 01. The AI Voice Buyer's Guide
- 02. What to look for, what to refuse to pay for, and the seven questions every vendor dodges.
- 03. Why this exists
- 04. The five things that actually matter
- 05. The seven questions vendors dodge
- 06. Pricing reality
- 07. What we'd do in your shoes
- 08. How to reach us

The AI Voice Buyer's Guide

What to look for, what to refuse to pay for, and the seven questions every vendor dodges.

By the team at TrainYourAgent — Tampa Bay, FL · 2026 edition

Why this exists

The AI voice market in 2026 is loud. Every vendor pitches "natural-sounding agents" with the same five demos, the same soft-ball calls, and pricing pages built to confuse you. This guide is a 15-minute read. It will give you the framework to vet any voice-agent vendor — including us.

The five things that actually matter

1. Trained on YOUR data, not a generic prompt

A real agent ingests your scripts, pricing, FAQs, ticket history, and product docs. A demo built on a system prompt isn't training — it's a wrapper.

How to vet: Ask "what data sources is the agent trained on?" If the answer is "we use the latest LLMs," it's a wrapper. If the answer involves your CRM, your ticket history, your transcripts — that's training.

2. Native integrations to YOUR stack

Booking should write into the calendar you actually use. Tickets should land in the CRM you actually use. Not webhooks duct-taped together.

How to vet: Name the three tools you use most. If the vendor says "yes, we integrate via Zapier," that's duct tape. If they name the integration product they've built, ask to see it live.

3. Multi-model fallback

Anthropic, OpenAI, Google all have outages. A serious vendor routes between them. A weekend vendor goes down when their primary model has a bad afternoon.

How to vet: Ask "what happens if your LLM provider has an outage?" If they pause, the answer is "you go down."

4. Real production observability

You should see every call's transcript, cost, quality score, and outcome in a dashboard. Not "we send you a weekly email."

How to vet: Ask for a screenshot of the dashboard a current customer sees. If they can't show you, the dashboard doesn't exist.

5. Tuning, not just deployment

The agent on Day 1 is not the agent on Day 90. Real vendors have a tuning process — sampling calls weekly, updating prompts, rolling back changes that hurt outcomes.

How to vet: Ask "who at your company will tune my agent and how often?" Vague answers = no process.

The seven questions vendors dodge

1. What's your average per-call cost (not per-minute)? Per-minute pricing without per-call context is meaningless.
2. Show me a recording of one of your live agent calls. Not a demo — a real call from a real customer.
3. What's your SLA on uptime, and how do you compensate when you breach it? No SLA = no skin in the game.
4. Who owns the data on cancel day? Anything other than "you do, and we delete our copy in 30 days" is a red flag.
5. Do you train AI models on customer data? If yes, walk away.
6. What's the largest agent you've shipped to production? Vague answers = small portfolio.
7. Can I talk to one of your customers? If they refuse, the customers don't exist or aren't happy.

Pricing reality

The cheapest per-minute on the market in 2026 is around \$0.05-0.08. The most expensive is \$0.50+. The right number for you depends on conversion: a \$0.05 agent that books 20% of qualifying calls costs you more per booking than a \$0.20 agent that books 55%. Always ask for cost-per-outcome, not cost-per-minute.

A reasonable build fee in 2026 is \$3,500-\$10,000 depending on integration complexity. Anything under \$1,500 is an off-the-shelf wrapper. Anything over \$25,000 should come with named senior engineers and a multi-month delivery plan.

What we'd do in your shoes

Pick three vendors. Give each the same scenario (your top 5 most common call types). Get a cost-per-outcome quote from each. Ask for a 7-day live trial. The vendor that says yes to all three goes first.

How to reach us

Honest disclosure: we're TrainYourAgent. We wrote this guide because the AI voice market needs better-informed buyers, and better-informed buyers tend to pick us anyway. If you want a 30-minute scoping call: cal.com/trainyouragent/30min. We'll show you a build live for your specific use case, give you a price on the spot, and you decide.

— Alexander Mills, Founder

Tampa Bay, FL · linkedin.com/in/alexandermillsai